



Programming Massively Parallel Hardware – Optimising Tridag


Group Project
Department of Computer Science

Written by:

Morten Espensen & Niklas Høj & Mathias Svennson
dzt440 - 19/06/1991 & nwt762 - 24/07/1991 & tpx783 - din/mor/XX
October 31, 2014

Supervised by:

Cosmin E. Oancea



Contents

1	Transformations	2
1.1	Tridag Rewrite	2
1.2	Naive CUDA implementation	2
1.3	Propagating outer loop into rollback	2
1.4	Optimising coalesced memory access	2
2	Will it validate?	3
3	Results	4

Chapter 1

Transformations

Our initial approach to implementing a CUDA version in this project was just to get started with converting the different parts of the program into a CUDA ready format, thus simplifying parts and transforming parts which were not easily implemented in CUDA kernels.

1.1 Tridag Rewrite

1.2 Naive CUDA implementation

1.3 Propagating outer loop into rollback

1.4 Optimising coalesced memory access

Chapter 2

Will it validate?

Yes.

Chapter 3

Results

Data set size / Implementation	Small	Medium	Large
Sequential with flatten arrays	2162560 μs	5652265 μs	195845401 μs
OpenMP	183016 μs	241972 μs	9680948 μs
Naive CUDA	3956322 μs	3639421 μs	34980508 μs
Optimised CUDA	183016 μs	241972 μs	9680948 μs